


计算机科学与技术学院神经网络与深度学习课程实验报告

实验题目：实现各种 classifier		学号：201900130143
日期：9/30	班级：智能	姓名：吴家麒
Email：996362192@qq.com		
<p>实验目的：</p> <ul style="list-style-type: none">• understand the basic Image Classification pipeline and the data-driven approach (train/predict stages)• understand the train/val/test splits and the use of validation data for hyperparameter tuning.• develop proficiency in writing efficient vectorized code with numpy• implement and apply a k-Nearest Neighbor (kNN) classifier• implement and apply a Multiclass Support Vector Machine (SVM) classifier• implement and apply a Softmax classifier• implement and apply a Three layer neural network classifier• understand the differences and tradeoffs between these classifiers• get a basic understanding of performance improvements from using higher-level representations than raw pixels (e.g. color histograms, Histogram of Gradient (HOG) features)		
<p>实验软件和硬件环境：</p> Anaconda3+Jupyter Notebook		
<p>实验原理和方法：</p> <p>1. KNN classifier 实现</p> <p>加载数据集：</p> <hr/> <pre>Training data shape: (50000, 32, 32, 3) Training labels shape: (50000,) Test data shape: (10000, 32, 32, 3) Test labels shape: (10000,)</pre> <p>预览部分训练数据：</p>		



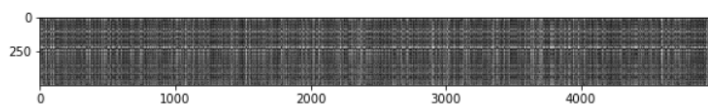
compute_distances_two_loops: 使用两重循环计算 distance matrix

```
# Open sdocs2019/classifiers/k_nearest_neighbor.py and implement
# compute_distances_two_loops.
```

```
# Test your implementation:
dists = classifier.compute_distances_two_loops(X_test)
print(dists.shape)
```

```
(500, 5000)
```

```
# We can visualize the distance matrix: each row is a single test example and
# its distances to training examples
plt.imshow(dists, interpolation='none')
plt.show()
```



分别用 k=1,

```
Got 137 / 500 correct => accuracy: 0.274000
```

和 k=5 计算预测精度:

```
Got 139 / 500 correct => accuracy: 0.278000
```

使用部分矢量化和一重循环加速 distance matrix 的计算:

```
# Now lets speed up distance matrix computation by using partial vectorization
# with one loop. Implement the function compute_distances_one_loop and run the
# code below:
```

```
dists_one = classifier.compute_distances_one_loop(X_test)
```

```
# To ensure that our vectorized implementation is correct, we make sure that it
# agrees with the naive implementation. There are many ways to decide whether
# two matrices are similar; one of the simplest is the Frobenius norm. In case
# you haven't seen it before, the Frobenius norm of two matrices is the square
# root of the squared sum of differences of all elements; in other words, reshape
# the matrices into vectors and compute the Euclidean distance between them.
```

```
difference = np.linalg.norm(dists - dists_one, ord='fro')
```

```
print('One loop difference was: %f' % (difference, ))
```

```
if difference < 0.001:
```

```
    print('Good! The distance matrices are the same')
```

```
else:
```

```
    print('Uh-oh! The distance matrices are different')
```

```
One loop difference was: 0.000000
```

```
Good! The distance matrices are the same
```

完全矢量化而不使用循环:

```
# Now implement the fully vectorized version inside compute_distances_no_loops
# and run the code
dists_two = classifier.compute_distances_no_loops(X_test)


# check that the distance matrix agrees with the one we computed before:
difference = np.linalg.norm(dists - dists_two, ord='fro')
print('No loop difference was: %f' % (difference, ))
if difference < 0.001:
    print('Good! The distance matrices are the same')
else:
    print('Uh-oh! The distance matrices are different')
```

No loop difference was: 0.000000
Good! The distance matrices are the same

比较上述几种方法的效率：

Two loop version took 31.478619 seconds
One loop version took 66.148429 seconds
No loop version took 0.213431 seconds

使用交叉验证确定超参的最佳值：



2. SVM classifier 实现

导入数据集：

```
Training data shape: (50000, 32, 32, 3)
Training labels shape: (50000,)
Test data shape: (10000, 32, 32, 3)
Test labels shape: (10000,)
```

分为训练集、验证集、训练集：

```
Train data shape: (49000, 32, 32, 3)
Train labels shape: (49000,)
Validation data shape: (1000, 32, 32, 3)
Validation labels shape: (1000,)
Test data shape: (1000, 32, 32, 3)
Test labels shape: (1000,)
```

再对数据进行一系列预处理后，调用 SVM classifier, 通过 naïve 和 vectorized 两

种方法进行实现，进行比较：


```
Naive loss: 9.120340e+00 computed in 0.123661s
Vectorized loss: 9.120340e+00 computed in 0.003963s
difference: 0.000000
```

```
Naive loss and gradient: computed in 0.132611s
Vectorized loss and gradient: computed in 0.002993s
difference: 0.000000
```

利用 SGD 方法最小化损失函数：

```
: # In the file linear_classifier.py, implement SGD in the function
# LinearClassifier.train() and then run it with the code below.
from sducs2019.classifiers import LinearSVM
svm = LinearSVM()
tic = time.time()
loss_hist = svm.train(X_train, y_train, learning_rate=1e-7, reg=2.5e4,
                      num_iters=1500, verbose=True)
toc = time.time()
print('That took %fs' % (toc - tic))
```

```
iteration 0 / 1500: loss 396.524657
iteration 100 / 1500: loss 236.428358
iteration 200 / 1500: loss 144.157705
iteration 300 / 1500: loss 88.272558
iteration 400 / 1500: loss 54.764083
iteration 500 / 1500: loss 35.191992
iteration 600 / 1500: loss 23.452266
iteration 700 / 1500: loss 16.118776
iteration 800 / 1500: loss 11.755828
iteration 900 / 1500: loss 9.386926
iteration 1000 / 1500: loss 7.177540
iteration 1100 / 1500: loss 6.915250
iteration 1200 / 1500: loss 5.918573
iteration 1300 / 1500: loss 5.579504
iteration 1400 / 1500: loss 5.331826
That took 6.926160s
```



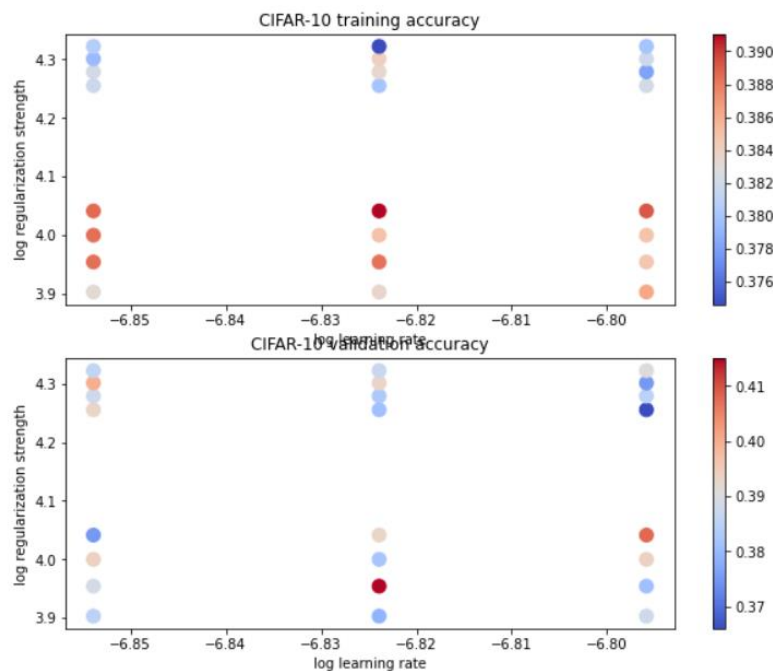
将模型用于验证集和测试集：

```
training accuracy: 0.381082
validation accuracy: 0.379000
```

交叉验证调参：

```
lr 1.600000e-07 reg 2.100000e+04 train accuracy: 0.379918 val accuracy: 0.391000
best validation accuracy achieved during cross-validation: 0.415000
```

可视化：



评估测试集最终的 accuracy：

```
: # Evaluate the best svm on test set
y_test_pred = best_svm.predict(X_test)
test_accuracy = np.mean(y_test == y_test_pred)
print('linear SVM on raw pixels final test set accuracy: %f' % test_accuracy)

linear SVM on raw pixels final test set accuracy: 0.362000
```

3. Softmax classifier 实现

导入数据集：

```
Train data shape: (49000, 3073)
Train labels shape: (49000,)
Validation data shape: (1000, 3073)
Validation labels shape: (1000,)
Test data shape: (1000, 3073)
Test labels shape: (1000,)
dev data shape: (500, 3073)
dev labels shape: (500,)
```

通过 naïve 和 vectorized 两种方法计算 Softmax Loss：

```
naive loss: 2.338143e+00 computed in 0.079806s
vectorized loss: 2.338143e+00 computed in 0.002995s
Loss difference: 0.000000
Gradient difference: 0.000000
```

结果相同但 Vectorized 更快。

交叉验证调整超参：

```
optimization time : 29.482837s
lr 1.000000e-07 reg 2.500000e+04 train accuracy: 0.251143 val accuracy: 0.251000
lr 1.000000e-07 reg 5.000000e+04 train accuracy: 0.241959 val accuracy: 0.243000
lr 5.000000e-07 reg 2.500000e+04 train accuracy: 0.314939 val accuracy: 0.315000
lr 5.000000e-07 reg 5.000000e+04 train accuracy: 0.311245 val accuracy: 0.301000
best validation accuracy achieved during cross-validation: 0.315000
```

计算测试集 accuracy:

```
# evaluate on test set
# Evaluate the best softmax on test set
y_test_pred = best_softmax.predict(X_test)
test_accuracy = np.mean(y_test == y_test_pred)
print('softmax on raw pixels final test set accuracy: %f' % (test_accuracy, ))

softmax on raw pixels final test set accuracy: 0.312000
```

4. 三层神经网络的实现

前向传播计算得分:

```
Your scores:
[[-0.03596154 -0.01613583 -0.00048556]
 [-0.09480192  0.20724618 -0.09763798]
 [-0.07015667  0.17081869 -0.07675745]
 [ 0.01473052  0.09321097 -0.0395178 ]
 [-0.05306489  0.04729807  0.01750587]]
```

correct scores:

```
Difference between your scores and correct scores:
4.807962381448306e-08
```

计算 loss:

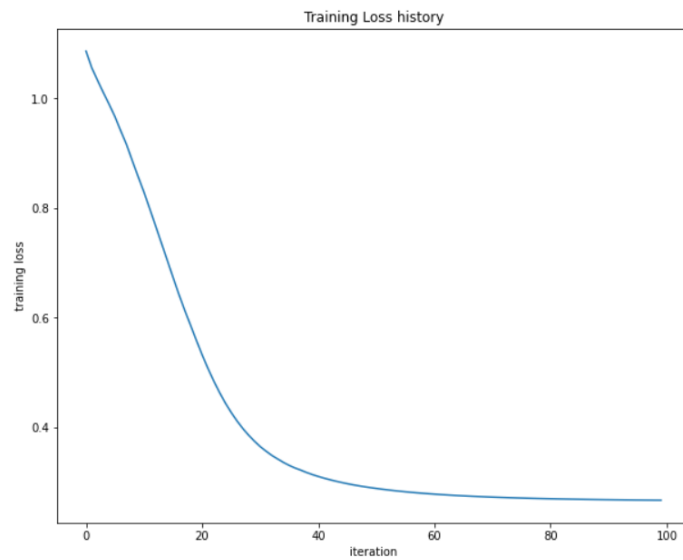
```
Difference between your loss and correct loss:
7.829292769656604e-13
```

通过 numeric gradient 验证反向传播的准确性, 得到 numeric 与 analytic 之间的差值:

```
W1 max relative error: 3.561318e-09
b1 max relative error: 7.670388e-09
W2 max relative error: 2.542856e-08
b2 max relative error: 4.590211e-10
W3 max relative error: 1.067978e-08
b3 max relative error: 8.037345e-11
```

使用 SGD 最小化损失:


Final training loss: 0.26699892976346745



导入数据集，训练神经网络：

```
iteration 0 / 2000: loss 2.302716
iteration 100 / 2000: loss 2.302714
iteration 200 / 2000: loss 2.302714
iteration 300 / 2000: loss 2.302715
iteration 400 / 2000: loss 2.302713
iteration 500 / 2000: loss 2.302713
iteration 600 / 2000: loss 2.302707
iteration 700 / 2000: loss 2.302713
iteration 800 / 2000: loss 2.302712
iteration 900 / 2000: loss 2.302710
iteration 1000 / 2000: loss 2.302714
iteration 1100 / 2000: loss 2.302712
iteration 1200 / 2000: loss 2.302706
iteration 1300 / 2000: loss 2.302708
iteration 1400 / 2000: loss 2.302704
iteration 1500 / 2000: loss 2.302706
iteration 1600 / 2000: loss 2.302708
iteration 1700 / 2000: loss 2.302706
iteration 1800 / 2000: loss 2.302708
iteration 1900 / 2000: loss 2.302713
Validation accuracy: 0.079
```

可视化损失下降过程和测试集验证集结果：



交叉验证调参：

```
iteration 0 / 2000: loss 2.306088
iteration 100 / 2000: loss 2.305985
iteration 200 / 2000: loss 2.305838
iteration 300 / 2000: loss 2.305759
iteration 400 / 2000: loss 2.305640
iteration 500 / 2000: loss 2.305550
iteration 600 / 2000: loss 2.305468
iteration 700 / 2000: loss 2.305386
iteration 800 / 2000: loss 2.305291
iteration 900 / 2000: loss 2.305152
iteration 1000 / 2000: loss 2.305219
iteration 1100 / 2000: loss 2.305097
iteration 1200 / 2000: loss 2.304984
iteration 1300 / 2000: loss 2.304973
iteration 1400 / 2000: loss 2.304945
iteration 1500 / 2000: loss 2.304850
iteration 1600 / 2000: loss 2.304852
iteration 1700 / 2000: loss 2.304814
iteration 1800 / 2000: loss 2.304741
iteration 1900 / 2000: loss 2.304811
0.079
```

Run on the test set:

```
: test_acc = (best_net.predict(X_test) == y_test).mean()
print('Test accuracy: ', test_acc)
```

Test accuracy: 0.1

5. High-Level Features

导入 CIFAR10 数据集，提取特征。

训练基于特征的 SVM，交叉验证后最终结果如下：

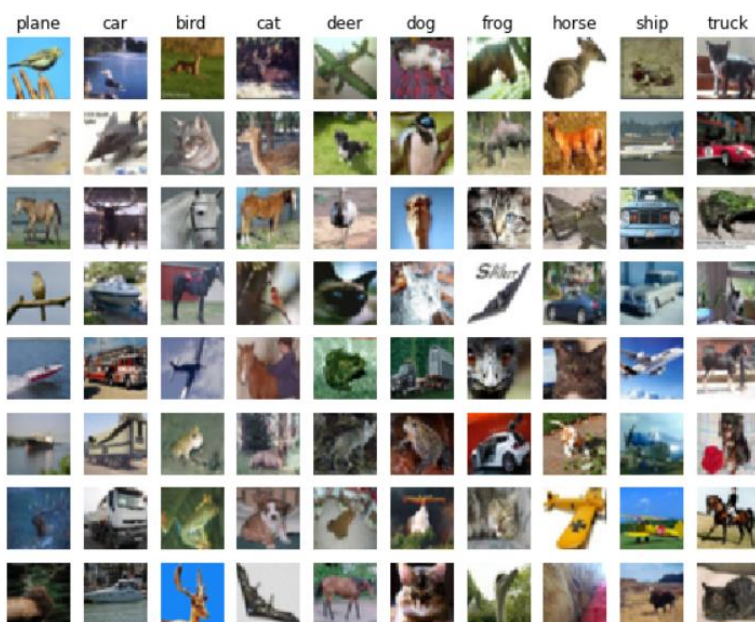

```
lr 1.000000e-07 reg 5.000000e+05 train accuracy: 0.405490 val accuracy: 0.401000
lr 1.000000e-07 reg 5.000000e+06 train accuracy: 0.359245 val accuracy: 0.365000
best validation accuracy achieved during cross-validation: 0.409000
```

测试集结果:

```
# Evaluate your trained SVM on the test set
y_test_pred = best_svm.predict(X_test_feats)
test_accuracy = np.mean(y_test == y_test_pred)
print(test_accuracy)
```

0.422

部分预测结果:



训练基于特征的神经网络,
设置学习率和正则化值:

```
learning_rates = [1e-4, 5e-3, 1e-3, 1e-2]
regularization_strengths = np.logspace(-4.3, -4.0, num=5) #[1e-4, 5e-4, 8e-4]
```

交叉验证后用于测试集得到结果:

```
: # Run your best neural net classifier on the test set. You should be able
# to get more than 55% accuracy.
```

```
test_acc = (best_net.predict(X_test_feats) == y_test).mean()
print(test_acc)
```

0.103

结论分析与体会：

1. 通过本次实验，回顾了几种常见的分类方法实现：KNN、SVM、Softmax，学习了构建三层神经网络学习器的方法，对于反向传播算法的内涵有了更深入的理解。
2. 对于学习模型的参数调整的过程也进行重新的熟悉，使用 SGD 最小化损失函数，通过交叉验证调整超参，最终运用于测试集查看模型训练的准确性。
3. 学习了对于各种训练模型的效率优化算法，对于提高模型学习效率有很大的帮助。
4. 通过对 CIFAR 数据集模型训练结果的直观体验，感受到了模型训练的实际效果与不足之处。

就实验过程中遇到和出现的问题，你是如何解决和处理的，自拟 1—3 道问答题：

1. 使用三层神经网络进行训练时，使用原始参数的效果不佳？
通过不断的调整学习率与正则化项进行优化，增加迭代次数，可以实现更好的训练效果。